



Eric Cuiffo, Nick Palumbo, Val A. Red, Aiser Sheikh
27 March 2013
198:416 Operating Systems, A02.

Fair Share Scheduler

- I. Introduction
- II. Areas of Change
- III. Implementation Details
- IV. Overhead
- V. Testing
- VI. Further improvement

I. Introduction:

Our group (incorporating Eric Cuiffo, Nick Palumbo, Val A. Red, and Aiser Sheikh) implemented a fair share scheduler in the Linux 2.4.24 kernel via a provided debian distribution. All user's processes will share an equal portion of CPU time in a way such that one user cannot use another user's equal portion. In addition, we implemented a system call that enables a user to request a double share to be allocated for its process. We called our system call `incshare`.

Our `patch_name.patch` file can be deployed in the Linux-2.4.24 directory via the following command

```
patch -p1 -i patch_name.patch
```

Data gathered from our implementation will be discussed in this report. We deployed on the Rutgers' Computer Science department's iLab Computers.

Invoking `incshare()`

We named our system call `incshare` instead of `myincshare`. It can be defined as followed:

```
12      #define __NR_incshare 259
13      _syscall0(int, incshare);
.
.
.
51      incshare()
```

II. Areas of Change:

Our modification of the kernel was mainly aimed at the files involved with the scheduler, with one addition of a system call, (`incshare`.) Overall, we added lines to the scheduling files to implement an algorithm that was a Fair Share Scheduler; essentially enabling multiple users logged into the machine to receive an equal allocation of CPU time for their processes.

The addition of the system call, `incshare()` also allowed users to give one of their processes a greater subdivision of CPU time out of their own allocation.

- As an overview, our `patch_name.patch` has edited the following files in the kernel:

```
arch/i386/kernel/entry.S
include/asm-i386/unistd.h
include/linux/sched.h
kernel/exit.c
kernel/fork.c
kernel/sched.c
kernel/sys.c           // Where incshare() is implemented and defined
Makefile
```

- To the process struct, we added `atomic_t share`.

This keeps track of how many times our system call `incshare` was called on a process; initializes as “1” by default.

- Added the initialization of `p->share` to 1 in `do_fork`
- On each repeat schedule, we counted number of processes (using `for_each_task`)

III. Implementation Details:

In the `task_struct` (which represents a process) there is a field called `counter`. In the original implementation of the Linux scheduler, this is initialized to the niceness value for a CPU-bound process. (For I/O bound processes, the previous counter value is taken into consideration and shifted right. For CPU-bound processes, the counter value reaches 0 before re-initialization.) As we are assuming all CPU-bound processes, the counter values will not be reinitialized until all the processes’ counters have reached 0 (technically only processes on the run queue, but we are assuming no blocking, so all processes should be on the run queue).

In order to implement the per-user scheduling algorithm defined in the assignment, we need to manipulate the value the counter is initialized to. Because the counter is decremented by one each time a process is scheduled, the ratio of initial counter values will correspond to the ratio of CPU time each process gets. (For example, if one process has a counter value of 5 and one has a counter value of 10, the second process will get twice as much CPU time before a re-initialization occurs).

Thus, in order to implement the desired scheduling, we initialized the counter for each process to the following calculation:

$$\text{counter} = \frac{\text{total number procs}}{\text{this user's number of procs}}$$

We can prove that this gives each user the same total amount of CPU time:

This equation ensures that each user’s collective counter total for their processes will be equal to the total number of processes in the system. If all their collective counter totals are equal, then each user is getting the same amount of process time. We can also observe that users with more processes will get less CPU time per process. In addition, each of a user’s processes (in the absence of a system call) will get the same CPU Time. This is congruent with our desired behavior. We also ensure that if this would result in a counter value of 0 for a process, we instead initialize it to 1.

Finally, in order to implement the system call, we used an `atomic_t` called `share` in the `task_struct`. When a process calls `incshare`, the `share` value increases from 1 to 2. In addition, the number of processes a user has is artificially increased by 1. We ensured that when a process that has called the `incshare` exits, the user’s number of processes is decreased by 2 to maintain consistency. The calculation for counter initialization is altered:

$$\text{counter} = \text{share} \cdot \frac{\text{total number procs}}{\text{this user's number of procs}}$$

It is clear that if no process calls `incshare`, the desired behavior is still observed. If a process calls `incshare`, it is clear that only that user’s processes have their counters altered and that process will get twice as much as the other process for that user.

*Defining the system call via **entry.S**, **unistd.h**, and **sys.c** :*

We first needed to define our system-call handling routine via the Entry.s file. We accomplished this by adding the following to linux/arch/i386/kernel/entry.S

```
666                                     .long SYMBOL_NAME(sys_incshare)
```

Adding a unistd.h entry and wrapper function

We then added an entry to include/asm-i386/unistd.h and proceeded to add our wrapper function:

```
260                                     #define __NR_incshare          259
.
.
.
379                                     static inline _syscall0(int,incshare)
```

Finally, we decided that we would add our system call through sys.c. In order to do this, we added the following lines to /kernel/sys.c :

```
1288                                     asmlinkage int sys incshare()
1289                                     {
.
.
.
1300                                     }
```

IV. Overhead:

The overhead generated by our changes includes the time it takes to calculate the total number of processes (which is equivalent to the time it takes to step through the processes in `for_each_task` and the time it takes to increment an integer). Also, the time it takes to initialize the `atomic_t` share when a process is created in `fork`. With regards to space, we are storing an additional `atomic_t` per process. Thus, we are using `sizeof(atomic_t)` additional space per process.

V. Testing:

For testing, we examined how our custom kernel and scheduler would run with multiple users and processes in the background to ensure that our Fair Share Scheduler actually succeeded in fairly allocating CPU time. At the same time, we also examined the functionality of our syscall `incshare()` to ensure that users could invoke it properly and have the process invoking `incshare()` receive double the allocation of CPU time compared with the other processes from the same user.

Our figures and additional details can be found on the following page:

Figure 1: Above displays our Fair Share Scheduler successfully allocating 25% CPU time to each of four users.

```

Elapsed time: 3.12 seconds

Debian-exim:      0% ( 0%, 0%)
(  exim4)  751    0% ( 0%, 0%)

User 9000:        25% ( 25%, 0%)
( cpuboun)  899   25% ( 25%, 0%)

User 9001:        25% ( 25%, 0%)
( cpuboun)  900   25% ( 25%, 0%)

User 9002:        25% ( 25%, 0%)
( cpuboun)  901   25% ( 25%, 0%)

User 9003:        25% ( 25%, 0%)
( cpuboun)  902   25% ( 25%, 0%)Quit
cs416:/home/user/new_ut/utilities# _

```

Compare the above image with the following, which is also a screen capture of our Fair Share Scheduler with a single user invoking `incshare()`, thus allowing one of the user's processes to receive a larger allocation than other processes from the same user:

Figure 2: Note how User 9000's processes receive 16% and 8% respectively as process 721 invokes `incshare()`. User 9001 does not invoke `incshare` at all, however, and thus both 9001's processes 723 and 724 both receive 12.5% allocations.

```

Elapsed time: 3.33 seconds

Debian-exim:      0% ( 0%, 0%)
(  exim4)  667    0% ( 0%, 0%)

User 9000:        24% ( 24%, 0%)
( cpuboun)  721   16% ( 16%, 0%)
( cpuboun)  722    8% (  8%, 0%)

User 9001:        25% ( 25%, 0%)
( cpuboun)  723   12% ( 12%, 0%)
( cpuboun)  724   12% ( 12%, 0%)

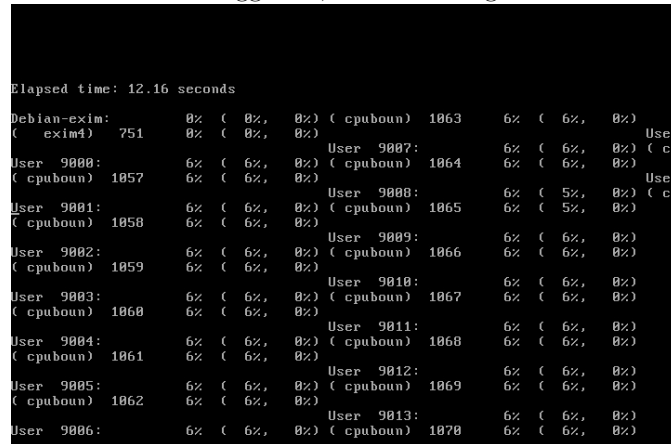
User 9002:        25% ( 25%, 0%)
( cpuboun)  725   25% ( 25%, 0%)

User 9003:        25% ( 25%, 0%)
( cpuboun)  726   25% ( 25%, 0%)_

```

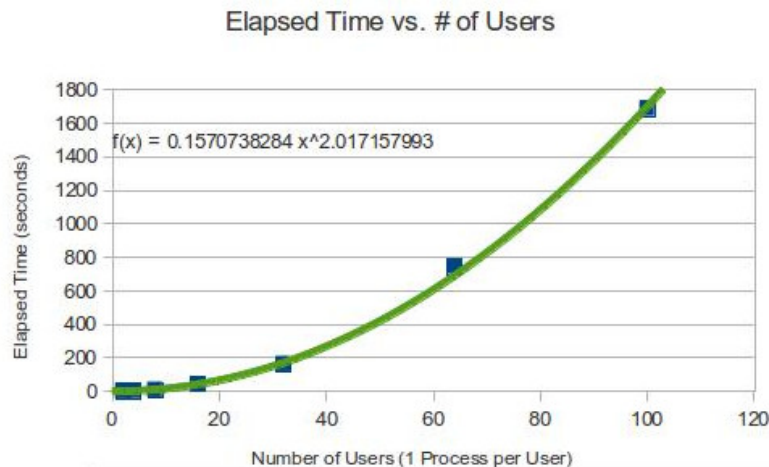
Our Fair Share Scheduler works no matter how many users are logged in and using the CPU. In the below image, 16 users are logged in and thus all are allocated 6% CPU time:

Figure 3: 16 users are logged in, each receiving 6% CPU allocation.



We went up to 100 users and were able to come up with the following graph, exemplifying an exponential, square regression:

Figure 4: This graph compares the time it takes to execute the last user's process out a range from 2 to 100 users.



At 100 simultaneous users, 1 process each, we successfully tested up to 100 simultaneous processes. Theoretically, our Fair Share Scheduler should even fairly schedule 10000 simultaneous processes; however, this would cause a starvation of resources from the user shell via a fork bomb, thus we were unable to observe our virtual machine under these conditions.

Overall, our tests were extensive and support what we know about Fair Share Scheduler and how it should be implemented to share resources between users.

VI. Further improvement:

Our `for_each_task` does not necessarily count every single process that exists in the system (we believe processes that do things w/r/t the kernel are not being counted, things such as `init`). However, our project is only concerned with the processes that we have control over and does not need to worry about the scheduling of those low-level processes. This may, however, cause a few percent discrepancy for CPU time in different processes. It may be worth at some point looking into a more accurate or quicker way to determine the number of processes in the system. In addition, as our implementation using a little extra space per process,

it may limit the total number of processes that can be in the system than usual—more resources as being used. Finally, in the presence of many I/O bound processes, our algorithm will fail. The counter values we used are only appropriate for CPU bound processes and will be inaccurate for I/O bound processes. However, our project operated under the assumption of only CPU-bound processes.